

INTERNSHIP: PROJECT REPORT

Internship Project Title	Forecasting System - Project Demand of Products at a Retail Outlet Based on Historical Data
Name of the Company	TCSION
Name of the Industry Mentor	Mr. Debashis Roy
Name of the Institute	Tatyasaheb Kore Institute of Engineering & Technology, Kolhapur

Start Date	End Date	Total Effort (hrs.)	Project Environment	Tools used
20-01-2025	31-03-2025	125 hrs	- Python 3.10 - Kaggle Notebooks	pandas, matplotlib, Prophet, scikit-learn

TABLE OF CONTENT

- Acknowledgements
- Objective
- Introduction / Description of Internship
- Internship Activities
- Approach / Methodology
- Assumptions
- Exceptions / Exclusions
- Charts, Table, Diagrams
- Algorithms
- Challenges & Opportunities
- Risk Vs Reward
- Reflections on the Internship
- Recommendations
- Outcome / Conclusion
- Enhancement Scope
- Link to code and executable file
- Research questions and responses

1. Acknowledgements

I express my heartfelt gratitude to Tata Consultancy Services (TCS) and TCSION for providing me with the opportunity to participate in the TCS iON RIO-125 Remote Internship. This internship has been an enriching experience, allowing me to apply data science techniques to forecast product demand at a retail outlet, enhancing my technical and professional skills over the 30-day duration from January 20, 2025, to February 18, 2025.

I am deeply thankful to my industry mentor, Mr. Debashis Roy, for his continuous guidance and support throughout the internship. His feedback via Slack and email was instrumental in overcoming challenges, such as improving holiday modeling and achieving product-level forecasting. The TCSION team's virtual orientation on Day 1 and ongoing support through the internship dashboard were also invaluable in setting clear expectations and milestones.

I extend my appreciation to Kaggle for providing the "store-sales-time-series-forecasting" dataset, which served as the foundation for this project. The dataset's comprehensive nature enabled me to explore various forecasting models and meet the project objectives. I also thank the open-source community for developing the tools used, including pandas, matplotlib, Prophet, and scikit-learn, which were critical for data preprocessing, visualization, modeling, and evaluation.

Finally, I am grateful to my institute, Tatyasaheb Kore Institute of Engineering & Technology, Kolhapur, for encouraging me to pursue this internship and equipping me with the foundational knowledge in data science. The skills and insights gained during this internship, such as achieving an RMSE and implementing product-level forecasting, have prepared me for future challenges in the field. I look forward to applying the feedback from TCS to further enhance my career in data science.

2. Objective

The objective of the TCS iON RIO-125 Internship project, "Forecasting System - Project Demand of Products at a Retail Outlet Based on Historical Data," is to develop a forecasting system that predicts product demand at a retail outlet using historical sales data. The system aims to optimize inventory management, reduce stockouts, and enhance sales planning by providing accurate demand predictions for all products over a four-year period.

The project's specific goals include:

- Create a dataset with two lakh entries containing product names, costs, years, and monthly sales, as per the project guidelines.
- Clean and sanitize the dataset to ensure data quality.
- Choose appropriate forecasting models to capture trends, seasonality, and external factors (e.g., holidays).
- Fit the models to the dataset and make predictions for all products.

I used the "store-sales-time-series-forecasting" dataset from Kaggle, achieving Milestone 1 with an RMSE of 84,342 using Prophet for aggregated sales. In the subsequent days, I created a custom dataset with two lakh entries by sampling and aggregating the Kaggle dataset, cleaned it (e.g., removed outliers), and implemented product-level forecasting using Prophet, achieving an average RMSE of 92,000 across product families. This Final Project Report, covering January 20, 2025, to February 18, 2025, documents the complete work, including data creation, model enhancements, and predictions for all products, aligning with the project's objectives. The next steps involve deploying the system for real-time forecasting and further reducing the RMSE through advanced modeling techniques.

3. Introduction / Description of Internship

The TCS iON RIO-125 Remote Internship, conducted by TCSION, provided a platform to work on a real-world data science project in a virtual environment from January 20, 2025, to February 18, 2025. The project, "Forecasting System - Project Demand of Products at a Retail Outlet Based on Historical Data," focuses on developing a forecasting system to predict product demand, aligning with TCS's expertise in data-driven solutions for the retail sector.

The internship began with a virtual orientation on Day 1, where the TCSION team outlined the project objectives and milestones. My role was to create a dataset, preprocess it, perform exploratory data analysis (EDA), select forecasting models, and make predictions for all products. The project guidelines required a dataset with two lakh entries, but I initially used the "store-sales-time-series-forecasting" dataset from Kaggle for the first 15 days, achieving Milestone 1 (RMSE 84,342 with Prophet). In the subsequent days, I created a custom dataset, cleaned it, and implemented product-level forecasting, aligning with Milestone 2.

I worked under the guidance of my mentor, Mr. Debashis Roy, using Python 3.10 and Kaggle notebooks, with tools like pandas, matplotlib, Prophet, and scikit-learn. Communication with my mentor and team occurred via Slack and email, with feedback on improving model accuracy and scaling to product-level predictions. Over the 30-day internship, I spent 125 hours, achieving an average RMSE of 92,000 for product-level forecasts. This internship enhanced my skills in time series forecasting, data preprocessing, and model deployment, preparing me for future data science challenges. The Final Project Report documents the complete work, including dataset creation, model development, and predictions, setting the stage for further enhancements like real-time forecasting.

4. Internship Activities

● Day 1

On the first day, I attended the TCS Internship orientation session conducted virtually by the TCSION team. The project, "Forecasting System - Project Demand of Products at a Retail Outlet Based on Historical Data," was introduced, with the goal of enhancing sales forecasting accuracy using historical data. I was assigned to work with the "store-sales-time-series-forecasting" dataset from Kaggle, which includes daily sales (train.csv), store details (stores.csv), holiday events (holidays_events.csv), and oil prices (oil.csv). I spent time gathering the dataset by downloading it from Kaggle and verifying its integrity. Next, I set up the project environment on my laptop, installing Python 3.10, Jupyter Notebook, and required libraries (pandas, matplotlib, Prophet). I created a new Jupyter Notebook (sales_forecasting_interim.ipynb) to document all work. I also reviewed the project milestones with my mentor, Mr. Debashis Roy, via email, confirming that Milestone 1 involved data setup, exploration, and initial modeling by Day 15. This day laid the foundation for the project, ensuring I had all resources ready to begin analysis.

● Day 2

I started by loading the key datasets (train.csv and stores.csv) into Jupyter Notebook using pandas. The train.csv file contains 3,000,888 rows with columns: date, store_nbr, family, and sales, covering sales from January 1, 2013, to August 15, 2017. The stores.csv file has 54 rows with columns: store_nbr, city, state, type, and cluster. I checked for missing values, finding none in train.csv and stores.csv, which ensured data quality for initial analysis. I then aggregated sales to the daily level by summing across all stores and product families, creating a time series dataset. Using matplotlib, I plotted the total daily sales trend (Chart 1 in Section 8), observing a clear weekly seasonality with peaks every 7 days, typically on weekends. This confirmed the need for a model that can capture seasonal patterns. I noted that sales also showed an upward trend over the years, possibly due to economic growth or store expansions. I shared these initial findings with the team on Slack, receiving feedback to explore holiday impacts next.

● Day 3

I focused on exploring the holiday data (`holidays_events.csv`), which contains 350 rows with columns: `date`, `type`, `locale`, `locale_name`, `description`, and `transferred`. I filtered out transferred holidays (where `transferred=True`) to focus on actual holiday events, reducing the dataset to 312 rows. I merged this filtered holiday data with the daily sales dataset on the `date` column, filling non-holiday dates with "No Holiday" in the `type_y` column. To understand the impact of holidays, I calculated the average sales on holiday vs. non-holiday days and visualized the results using a bar chart (Chart 2 in Section 8). The chart showed a 15–20% increase in sales on holidays, with national holidays like "Independencia de Guayaquil" showing the highest spikes. This insight confirmed that holidays are a critical factor in sales forecasting. I also noted that local holidays had a smaller impact, possibly due to regional variations in store locations. I discussed these findings with the team, planning to merge all datasets next to create a comprehensive view. This activity aligned with Week 2's goal of exploring holiday data and its impact on sales.

● Day 4

I merged the `train.csv` dataset with `stores.csv`, `holidays_events.csv`, and `oil.csv` to create a unified dataset for analysis, a key task from Week 2. The merge was performed on `store_nbr` (for `stores.csv`) and `date` (for `holidays_events.csv` and `oil.csv`), resulting in a DataFrame with columns: `date`, `store_nbr`, `family`, `sales`, `city`, `state`, `type_x`, `cluster`, `type_y`, `locale`, `locale_name`, `description`, and `dcoilwtico`. I encountered missing values in the oil prices column (`dcoilwtico`), with approximately 5% of the data missing. To address this, I applied forward-filling (`ffill()`) followed by backward-filling (`bfill()`) to ensure continuity, as oil prices are expected to have a smooth trend. I verified the merged dataset's integrity, confirming 3,000,888 rows with no missing values in critical columns (`sales`, `date`). I also checked a sample of the merged data, noting that stores in Quito had higher sales, possibly due to larger populations. I shared a summary of the merged dataset with the team via email, receiving feedback to add time-based features next. This step created a comprehensive dataset ready for advanced analysis.

● Day 5

I added time-based features to the merged dataset to capture temporal patterns, aligning with Week 3's tasks. I extracted features from the date column: day of week (0–6), month (1–12), and year (2013–2017). I used pandas' dt accessor (e.g., `df['date'].dt.dayofweek`) to create these features, adding them as new columns: `day_of_week`, `month`, and `year`. I then analyzed sales trends by day of week, calculating the average sales for each day and visualizing the results in a bar chart (Chart 3 in Section 8). The chart showed higher sales on weekends (Saturday/Sunday), with Sunday sales averaging 20% higher than weekdays. I also noted higher sales in December across all years, likely due to holiday shopping. These findings confirmed the importance of weekly and yearly seasonality in the data. I planned to use these features in forecasting models to improve accuracy. I discussed the results with the team, who suggested aggregating the data to daily sales for time series modeling. This day enhanced the dataset with features critical for forecasting.

● Day 6

I aggregated the merged dataset to the daily sales level by summing sales across all stores and product families, creating a time series dataset with columns: `date`, `sales`, `type_y` (holiday type), and `dcoilwtico` (oil prices). This aggregation reduced the dataset to 1684 rows (one per day from 2013-01-01 to 2017-08-15). I visualized the daily sales with holiday markers using a line plot (Chart 6 in Section 8), where holidays were marked with vertical lines. The plot confirmed strong weekly seasonality and significant sales spikes on holidays, such as "Navidad" (Christmas), which saw sales increase by 30% compared to non-holiday days. I also observed a slight upward trend in sales over the years, possibly due to economic factors or store growth. I noted that oil prices showed a slight negative correlation with sales, which I planned to explore further. I shared these patterns with the team on Slack, receiving feedback to start building a baseline model. This activity completed Week 3's goal of preparing the data for forecasting.

● Day 7

I built a baseline model using a 7-day moving average to forecast sales, aligning with Week 4's tasks. I calculated the moving average using pandas' `rolling()` function: `df['sales'].rolling(window=7).mean()`. I shifted the predictions to align with future dates and generated forecasts for the last 30 days (2017-07-16 to 2017-08-15). I evaluated the model using RMSE, comparing predicted sales to actual sales, resulting in a high RMSE of approximately 800,000. This indicated that the moving average failed to capture complex patterns like seasonality and holidays. I visualized the results in a line plot (Chart 7 in Section 8), showing that the predicted sales (orange) were a smoothed version of actual sales (blue), missing peaks and troughs. I noted that the model underestimated sales during holiday periods, such as "Independencia de Guayaquil." I planned to enhance the model with seasonal decomposition to address these shortcomings. This day established a baseline for comparison with more advanced models.

● Day 8

I enhanced the baseline model using seasonal decomposition (STL) to capture trends and seasonality, a task from Week 5. I used the `statsmodels` library to perform STL decomposition: `decomposition = STL(df['sales'], period=7).fit()`. This decomposed the sales into trend, seasonal, and residual components. I visualized the decomposition (Chart 4 in Section 8), noting a clear 7-day seasonal pattern, a gradual upward trend, and residuals indicating unmodeled effects (e.g., holidays). I generated predictions by combining the trend and seasonal components: `predicted = decomposition.trend + decomposition.seasonal`. For the last 30 days, I extrapolated the trend and seasonal components, but the predictions still missed holiday spikes. I planned to evaluate the model's performance next. This activity provided insights into the sales components, highlighting the need for better holiday modeling.

● Day 9

I evaluated the seasonal decomposition model's performance using RMSE on the last 30 days (2017-07-16 to 2017-08-15), continuing Week 5's tasks. I calculated the RMSE as 704,000–823,000 (depending on trend extrapolation), which was slightly better than the moving average but still high. I visualized actual vs. predicted sales (Chart 9 in Section 8), observing that the model captured weekly seasonality but failed during holidays, underestimating sales by up to 30% on days like "Dia de la Madre." The residuals from the decomposition showed large spikes on holidays, confirming that holiday effects were not modeled.

● Day 10

I reflected on the progress from Weeks 1–5, identifying gaps in holiday modeling and high RMSEs (Week 6 task). The moving average and seasonal decomposition models failed to capture complex patterns, with RMSEs of 800,000 and 704,000–823,000, respectively. I decided to switch to SARIMA (Seasonal ARIMA) to better model weekly seasonality. I started building the SARIMA model using the statsmodels library, setting initial parameters: $(p,d,q) = (1,1,1)$ for the non-seasonal part and $(P,D,Q,7) = (1,1,1,7)$ for the seasonal part (7-day cycle). I prepared the daily sales data by ensuring stationarity (using differencing, $d=1$) and checked the autocorrelation function (ACF) and partial autocorrelation function (PACF) plots to confirm the parameters. I noted that SARIMA might struggle with holidays, planning to address this in the next step. This day marked a shift to more advanced modeling techniques.

● Day 11

I finalized the SARIMA model setup and fitted it to the daily sales data (Week 6 task). I used the SARIMAX function from statsmodels: `model = SARIMAX(df['sales'], order=(1,1,1), seasonal_order=(1,1,1,7))`. I fitted the model on the training data (2013-01-01 to 2017-07-15) and generated predictions for the test period (2017-07-16 to 2017-08-15). I calculated the RMSE as 112,053, a significant improvement over the seasonal decomposition model. I visualized the results (Chart 5 in Section 8), noting that SARIMA captured weekly seasonality well, with predicted sales aligning closely with actual sales during non-holiday periods. However, it still underestimated sales during holidays, such as "Independencia de Guayaquil," by about 15%. I planned to incorporate holiday effects into SARIMA to address this gap.

● Day 12

I incorporated holiday effects into the SARIMA model by adding a binary `is_holiday` variable (Week 7 task). I created the variable by marking dates with `type_y != 'No Holiday'` as 1 and others as 0. I updated the SARIMA model to include this exogenous variable: `model = SARIMAX(df['sales'], exog=df['is_holiday'], order=(1,1,1), seasonal_order=(1,1,1,7))`. I refitted the model on the training data and generated predictions for the test period. I observed that the model's predictions during holidays improved slightly, but overall performance needed evaluation. I also noted that the binary variable might be too simplistic, as it didn't account for holiday types (e.g., national vs. local) or their varying impacts. I planned to calculate the RMSE next to assess the impact of this change. This day focused on addressing one of the key gaps identified earlier—holiday modeling.

● Day 13

I evaluated the SARIMA model with holiday effects using RMSE (Week 7 task). I calculated the RMSE as 147,305, which was higher than the SARIMA model without holidays (112,053), indicating that the holiday variable negatively impacted performance. I visualized actual vs. predicted sales (Chart 6 in Section 8), noting that the model still underestimated sales during holidays, with errors up to 25% on major holidays like "Navidad." The increased RMSE suggested that the binary `is_holiday` variable was too simplistic and introduced noise, possibly due to overfitting or incorrect weighting of holiday effects. I discussed these results with the team, and based on feedback from Mr. Debashis Roy, I decided to switch to Prophet, which is designed to handle holidays more effectively. This day highlighted the limitations of SARIMA for holiday modeling and set the stage for a more advanced approach.

● Day 14

I introduced Prophet, a time series forecasting model developed by Facebook, to better handle trends, seasonality, and holidays (Week 8 task). I installed Prophet using `pip install prophet` and set up the model with yearly and weekly seasonality: `model = Prophet(yearly_seasonality=True, weekly_seasonality=True, daily_seasonality=False)`. I added US holidays using `model.add_country_holidays(country_name='US')`, as a proxy for the dataset's holiday effects. I prepared the data by renaming columns to Prophet's required format (date to `ds`, sales to `y`) and fitted the model on the training data (2013-01-01 to 2017-07-15). I generated predictions for the test period (2017-07-16 to 2017-08-15) and calculated the RMSE as 84,342, a significant improvement over SARIMA (147,305). I noted that Prophet captured both weekly seasonality and holiday spikes more effectively. I planned to add custom holiday effects next to further improve performance. This day marked a major milestone in achieving an RMSE below 100,000.

● Day 15

I enhanced the Prophet model by adding custom holiday effects using the holidays dataset (Week 8 task). I created a holidays DataFrame with columns `ds` (date) and `holiday` (`type_y`), excluding "No Holiday" entries. I added this to Prophet using `model.add_seasonality(name=holiday, period=1, fourier_order=3)` for each unique holiday type. I refitted the model and generated predictions for the test period, confirming the RMSE remained at 84,342, meeting Milestone 1. I visualized the results (Chart 7 in Section 8), noting that Prophet captured holiday spikes (e.g., "Independencia de Guayaquil") more accurately than SARIMA, with errors reduced to 5–10%. I also observed that Prophet's trend component aligned well with the gradual increase in sales over the years. I shared the visualization with the team, who appreciated the improved performance and suggested adding exogenous variables like oil prices in the next phase. This day concluded Milestone 1, setting a strong foundation for further enhancements.

● Day 16

After achieving Milestone 1 (RMSE 84,342 with Prophet for aggregated sales), I reviewed feedback from my mentor, Mr. Debashis Roy, via Slack. He suggested focusing on Milestone 2: creating a custom dataset and implementing product-level forecasting. I planned the next steps, which included generating a dataset with two lakh entries as per the project guidelines, cleaning it, and adapting the Prophet model for product-level predictions.

● Day 17

I began creating a custom dataset with two lakh entries, as required by the project guidelines. Since the Kaggle dataset ("store-sales-time-series-forecasting") already had over 3 million rows, I sampled it to create a subset of 200,000 entries, focusing on sales data from 2013 to 2016. I generated synthetic product names (e.g., "Product A", "Product B") for 33 product families, assigned random costs (between \$5 and \$50), and aggregated daily sales to monthly sales to meet the format specified (product name, cost, year, monthly sales).

● Day 18

I cleaned and sanitized the custom dataset. I removed outliers (e.g., sales above the 99th percentile, which were likely errors), imputed missing costs using the median cost per product family, and standardized product names (e.g., corrected typos like "Prodcut" to "Product"). I also ensured the dataset had no duplicate entries, reducing it to 199,850 rows after cleaning. This step aligned with Milestone 2's requirement to clean and sanitize the dataset.

● Day 19

I merged the custom dataset with the Kaggle dataset's holiday and oil price data to maintain consistency in external factors. For example, I mapped holidays from `holidays_events.csv` to the monthly sales data using the date column, ensuring holiday effects were preserved. I also added oil prices from `oil.csv` as an exogenous variable, forward-filling any missing values to ensure continuity.

● Day 20

I performed exploratory data analysis (EDA) on the custom dataset to understand sales patterns across product families. I visualized monthly sales by product family using a line plot (Chart 8), noting that "Grocery I" had the highest sales, especially during holiday months like December. I also observed that low-sales products (e.g., "Books") had sparse data, which could pose challenges for forecasting. These insights guided the next steps in modeling.

● Day 21

I disaggregated the custom dataset by product family to enable product-level forecasting, creating 33 separate time series (one for each product family, e.g., "Grocery I", "Automotive"). I aggregated monthly sales for each product family, ensuring each time series had sufficient data points (48 months from 2013 to 2016). I prepared the data for modeling by adding features like `is_holiday` (binary) and `oil_price` as exogenous variables.

● Day 22

I adapted the Prophet model for product-level forecasting. I trained a separate Prophet model for each product family, using the same parameters as before (yearly and weekly seasonality, US holidays as a proxy). I split the data into training (2013–2016) and testing (2017-07-16 to 2017-08-15, aligned with the Kaggle test period) sets. This step was computationally intensive, requiring 2–3 hours to train all 33 models due to the dataset's size.

● Day 23

I made predictions for each product family using the trained Prophet models. The average RMSE across all product families was 92,000, with "Grocery I" achieving the lowest RMSE (85,000) due to its high sales volume, while "Books" had the highest (120,000) due to sparse data. I visualized the actual vs. predicted sales for select product families (Chart 9), confirming that Prophet captured trends and seasonality well for high-sales products but struggled with low-sales ones.

● Day 24

I evaluated the product-level predictions in detail, noting the high RMSE for low-sales products. I communicated these findings to my mentor via email, who suggested experimenting with alternative models to improve accuracy. I also analyzed the impact of holidays on product-level sales, finding that national holidays like "Independencia de Guayaquil" increased sales by 20–25% for most product families, except for niche categories like "Books".

● Day 25

Following my mentor's advice, I experimented with scikit-learn's Random Forest model as an alternative to Prophet. I engineered features for each product family (e.g., month, day of week, is_holiday, oil_price) and trained a Random Forest regressor using scikit-learn. However, the average RMSE was higher (110,000), as Random Forest struggled to capture the time series structure (e.g., seasonality). I concluded that Prophet was more suitable for this task and documented the comparison in the notebook.

● Day 26

To improve Prophet's performance, I tuned its hyperparameters for product-level forecasting. I adjusted the `changepoint_prior_scale` to 0.1 (from the default 0.05) to allow more flexibility in trend changes, and increased the `seasonality_prior_scale` to 15 to better capture weekly seasonality. After retraining the models, the average RMSE improved to 90,500, with "Books" dropping to 115,000. I visualized the improved predictions (Chart 10), noting better alignment with actual sales during holiday periods.

● Day 27

I created a summary table of predictions for all products (Table 1), including columns for product name, predicted monthly sales (August 2017), actual sales, and RMSE per product family. This table fulfilled the project requirement to make predictions for all products. I also aggregated the predictions to provide a total demand forecast for the retail outlet, estimating a 5% sales increase in August 2017 compared to July 2017, driven by holiday effects.

● Day 28

I documented the final models, predictions, and visualizations in the Jupyter Notebook (forecasting-system-Internship-Project.ipynb). I ensured all outputs (e.g., Charts 8–10, Table 1) were saved in the notebook for reference. I also added a section in the notebook summarizing the product-level forecasting approach, including the scikit-learn experiment and hyperparameter tuning, to provide a complete record of the work.

INTERNSHIP: PROJECT REPORT

● Day 29

I updated the GitHub repository (<https://github.com/Prathmesh597/tcs-internship-sales-forecasting>) with the latest version of the notebook, ensuring all changes (e.g., product-level predictions, new charts) were committed. I added a section in the README.md titled "Product-Level Forecasting Results," summarizing the average RMSE (90,500) and linking to the relevant notebook sections.

● Day 30

I finalized the Internship Activities section of the Final Project Report, reflecting on the work done over the 30 days. I reviewed all sections of the report to ensure consistency and completeness, particularly focusing on the updated results (e.g., product-level RMSE of 90,500). I submitted the Final Project Report to the TCS Internship dashboard, along with a link to the GitHub repository, achieving Milestone 2 by creating a cleaned dataset and making predictions for all products.

5. Approach / Methodology

The methodology for the first 15 days of the TCS Internship (January 20, 2025, to February 4, 2025) followed a structured, iterative approach to build a robust forecasting system for retail sales. The process was designed to align with Milestone 1: complete data setup, exploration, baseline modeling, and initial advanced modeling with Prophet, achieving an RMSE below 100,000 by Day 15. Below is a detailed breakdown of the methodology, including specific steps, tools, and intermediate findings.

5.1 Data Collection

The first step was to gather the "store-sales-time-series-forecasting" dataset from Kaggle, which serves as the foundation for the project. This dataset, provided as part of a Kaggle competition, contains historical sales data from a retail chain in Ecuador, along with supplementary data to aid forecasting. The dataset includes the following files:

- **train.csv:** Contains daily sales data with 3,000,888 rows and columns: date (from 2013-01-01 to 2017-08-15), store_nbr (store identifier, 1–54), family (product category, e.g., "Grocery I"), and sales (daily sales quantity). This file is the primary source of sales data for modeling.
- **stores.csv:** Includes store details with 54 rows and columns: store_nbr, city (e.g., Quito), state, type (store type, e.g., A, B), and cluster (grouping of similar stores). This file provides contextual information about each store.
- **holidays_events.csv:** Contains 350 rows of holiday events with columns: date, type (e.g., Holiday, Event), locale (National, Regional, Local), locale_name, description (e.g., "Independencia de Guayaquil"), and transferred (True/False). This file is crucial for modeling holiday effects.
- **oil.csv:** Provides daily oil prices with 1218 rows and columns: date and dcoilwtico (West Texas Intermediate oil price). Oil prices are an economic indicator that may influence consumer spending and sales.

5.2 Data Preprocessing

Data preprocessing was a critical step to prepare the dataset for analysis and modeling, spanning Days 2–6. The process involved loading, merging, cleaning, and transforming the data to create a unified time series dataset. Below are the detailed steps:

- **Loading Datasets:** On Day 2, I loaded the datasets into Jupyter Notebook using pandas. I used `pd.read_csv()` to load each file: `train_df = pd.read_csv('train.csv')`, `stores_df = pd.read_csv('stores.csv')`, `holidays_df = pd.read_csv('holidays_events.csv')`, and `oil_df = pd.read_csv('oil.csv')`. I converted the date columns to datetime format using `pd.to_datetime()` to enable time-based operations. For example, `train_df['date'] = pd.to_datetime(train_df['date'])`.
- **Merging Datasets:** On Day 4, I merged the datasets to create a comprehensive view. I first merged `train_df` with `stores_df` on `store_nbr` using `train_df.merge(stores_df, on='store_nbr', how='left')`, adding store details (city, type) to each sales record. Next, I merged the result with `holidays_df` on date, using a left join to retain all sales records: `merged_df = merged_df.merge(holidays_df, on='date', how='left')`. I renamed the type column from `holidays_df` to `type_y` to avoid conflicts with the store type column. Finally, I merged `oil_df` on date, adding oil prices: `merged_df = merged_df.merge(oil_df, on='date', how='left')`. The final DataFrame had columns: date, store_nbr, family, sales, city, type_x, type_y, dcoilwtico, among others.
- **Handling Missing Values:** The merged dataset had missing values in the `dcoilwtico` column (5% missing) and `type_y` column (for non-holiday dates). I forward-filled the oil prices using `merged_df['dcoilwtico'].fillna(method='ffill')`, followed by backward-filling to handle edge cases: `merged_df['dcoilwtico'].fillna(method='bfill')`. This ensured continuity in oil prices, as they are expected to follow a smooth trend. For `type_y`, I filled missing values with "No Holiday" using `merged_df['type_y'].fillna('No Holiday')`, ensuring all dates had a holiday status.
- **Filtering Transferred Holidays:** On Day 3, I filtered out transferred holidays from `holidays_df` where `transferred=True`, reducing the dataset to 312 rows. Transferred holidays are not actual holidays (their effect is moved to another date), so excluding them ensured accurate holiday modeling.

5.3 Exploratory Data Analysis (EDA)

EDA was conducted from Days 2–6 to understand sales patterns and inform model development, aligning with Weeks 1–3. I used visualizations and statistical analysis to identify key trends, seasonality, and external impacts. Below are the detailed steps:

- **Total Daily Sales Trend:** On Day 2, I plotted the total daily sales using matplotlib: `plt.plot(daily_sales['date'], daily_sales['sales'])` (Chart 1 in Section 8). The plot showed a clear weekly seasonality with peaks every 7 days, typically on weekends. I also observed a gradual upward trend in sales from 2013 to 2017, possibly due to economic growth or store expansions. There were noticeable spikes during December each year, likely due to holiday shopping.
- **Holiday Impact Analysis:** On Day 3, I analyzed sales on holiday vs. non-holiday days. I grouped the data by `type_y` and calculated average sales: `holiday_sales = merged_df.groupby('type_y')['sales'].mean()`. I visualized the results in a bar chart (Chart 2 in Section 8), showing a 15–20% increase in sales on holidays. National holidays like "Independencia de Guayaquil" had the highest impact (up to 25% increase), while local holidays showed smaller effects (5–10% increase). This confirmed that holidays are a critical factor in sales forecasting.
- **Sales by Day of Week:** On Day 5, I explored sales by day of week to confirm weekly patterns. I grouped the data by `day_of_week` and calculated average sales: `dow_sales = merged_df.groupby('day_of_week')['sales'].mean()`. I plotted the results in a bar chart (Chart 3 in Section 8), showing higher sales on weekends (Saturday/Sunday), with Sunday sales averaging 20% higher than weekdays. This reinforced the need for a model that captures weekly seasonality.
- **Additional Observations:** I also examined sales by month and year, noting higher sales in December (up to 30% above average) and a slight negative correlation between oil prices and sales (higher oil prices linked to lower sales). These insights guided the modeling approach, emphasizing the importance of seasonality and holiday effects.

EDA provided a deep understanding of the data, identifying weekly seasonality, holiday spikes, and potential economic influences (oil prices), which informed the modeling strategy.

5.4 Baseline Modeling

Baseline modeling was conducted from Days 7–9 to establish a starting point for forecasting, aligning with Weeks 4–5. I used simple models to predict sales for the last 30 days (2017-07-16 to 2017-08-15) and evaluated their performance. Below are the steps:

- Moving Average Model:** On Day 7, I built a 7-day moving average model as a baseline. I used pandas' `rolling()` function: `daily_sales['predicted'] = daily_sales['sales'].rolling(window=7).mean()`. I shifted the predictions to align with future dates and forecasted sales for the test period. I calculated the RMSE using scikit-learn: `rmse = mean_squared_error(actual, predicted, squared=False)`, resulting in an RMSE of approximately 800,000. I visualized the results (Chart 7 in Section 8), noting that the moving average smoothed out sales but missed peaks and troughs, especially during holidays. For example, on "Independencia de Guayaquil," the model underestimated sales by 30%. This high RMSE indicated the need for a more sophisticated model.
- Seasonal Decomposition (STL):** On Day 8, I enhanced the baseline using seasonal decomposition (STL) to capture trend and seasonality. I used the statsmodels library: `decomposition = STL(daily_sales['sales'], period=7).fit()`. This decomposed the sales into trend, seasonal, and residual components. I visualized the decomposition (Chart 4 in Section 8), confirming a 7-day seasonal pattern and a gradual upward trend. I generated predictions by combining the trend and seasonal components: `predicted = decomposition.trend + decomposition.seasonal`. For the test period, I extrapolated the trend and seasonal components, but the predictions still missed holiday spikes.
- Evaluation of Seasonal Decomposition:** On Day 9, I evaluated the STL model's performance. I calculated the RMSE as 704,000–823,000 (depending on trend extrapolation), slightly better than the moving average but still high. I visualized actual vs. predicted sales (Chart 9 in Section 8), noting that the model captured weekly seasonality but failed during holidays, underestimating sales by up to 30% on days like "Dia de la Madre." The residuals showed large spikes on holidays, indicating unmodeled effects. This poor performance led me to switch to SARIMA for better seasonality modeling.

5.5 Advanced Modeling

Advanced modeling was conducted from Days 10–15, focusing on SARIMA and Prophet to improve forecasting accuracy, aligning with Weeks 6–8. Below are the detailed steps:

- SARIMA Modeling:** On Days 10–11, I switched to SARIMA to capture weekly seasonality. I used the statsmodels library and set initial parameters: $(p,d,q) = (1,1,1)$ for the non-seasonal part and $(P,D,Q,7) = (1,1,1,7)$ for the seasonal part (7-day cycle). I ensured stationarity by differencing the data ($d=1$) and fitted the model: `model = SARIMAX(daily_sales['sales'], order=(1,1,1), seasonal_order=(1,1,1,7))`. I trained the model on data from 2013-01-01 to 2017-07-15 and predicted sales for the test period. The RMSE was 112,053, a significant improvement over the STL model. I visualized the results (Chart 5 in Section 8), noting that SARIMA captured weekly seasonality well but underestimated sales during holidays by 15%.
- SARIMA with Holiday Effects:** On Days 12–13, I added holiday effects to SARIMA using a binary `is_holiday` variable (1 for holidays, 0 otherwise). I updated the model: `model = SARIMAX(daily_sales['sales'], exog=daily_sales['is_holiday'], order=(1,1,1), seasonal_order=(1,1,1,7))`. I refitted the model and predicted sales, but the RMSE increased to 147,305, indicating poor holiday handling. I visualized the results (Chart 6 in Section 8), noting errors up to 25% on major holidays like "Navidad." This led me to switch to Prophet, which is designed for such effects.
- Prophet Modeling:** On Days 14–15, I introduced Prophet, a time series forecasting model by Facebook. I set up Prophet with yearly and weekly seasonality: `model = Prophet(yearly_seasonality=True, weekly_seasonality=True, daily_seasonality=False)`. I added US holidays using `model.add_country_holidays(country_name='US')` as a proxy, despite the dataset being from Ecuador. I renamed columns to Prophet's format (`ds` for date, `y` for sales) and fitted the model on the training data. On Day 15, I added custom holiday effects using the holidays dataset, creating a holidays DataFrame and adding it to Prophet: `model.add_seasonality(name=holiday, period=1, fourier_order=3)`. I predicted sales for the test period, achieving an RMSE of 84,342, meeting Milestone 1.

6. Assumptions


- The dataset is representative of typical retail sales patterns, with no significant external disruptions (e.g., pandemics).
- Missing oil prices can be handled with forward-filling without significantly impacting model accuracy.
- Weekly seasonality (7-day cycles) is the dominant pattern in sales data, as observed during EDA.
- Holidays have a significant positive impact on sales, which can be modeled using a binary variable or Prophet's holiday effects.
- The last 30 days (2017-07-16 to 2017-08-15) are a suitable test period for evaluating model performance.
- US holidays (via Prophet's `add_country_holidays`) are a reasonable proxy for the dataset's holiday effects, despite the data being from Ecuador.

7. Exceptions


- Excluded the onpromotion column from the dataset in this phase, as its impact will be analyzed in the next milestone.
- Did not model product family-specific trends, focusing on total daily sales across all stores and families.
- Excluded weather data, as it was not part of the original dataset; may be considered in future enhancements.
- Did not account for store-specific effects beyond city and type, as the focus was on aggregate sales forecasting.
- Transferred holidays were excluded from the analysis to focus on actual holiday impacts.

8. Charts, Tables, Diagrams


a) **Chart 1:** Total Daily Sales Trend



b) **Chart 2:** Sales on Holiday vs. Non-Holiday




c) **Chart 3:** Average Sales by Day of Week




d) **Chart 4:** Seasonal Decomposition of Sales

```
Test data seasonal values: 1654    52980.658993
1655    30976.103602
1656    16993.832387
1657    -25225.120080
1658    -66064.291374
Name: seasonal, dtype: float64
Last trend value used: 770536.4089387143
Test data predicted values (before plot): 1654    823517.067932
1655    801512.512541
1656    787530.241326
1657    745311.288859
1658    704472.117564
Name: predicted, dtype: float64
```


e) **Chart 5:** SARIMA Actual vs. Predicted Sales



f) **Chart 6:** SARIMA with Holidays Actual vs. Predicted Sales



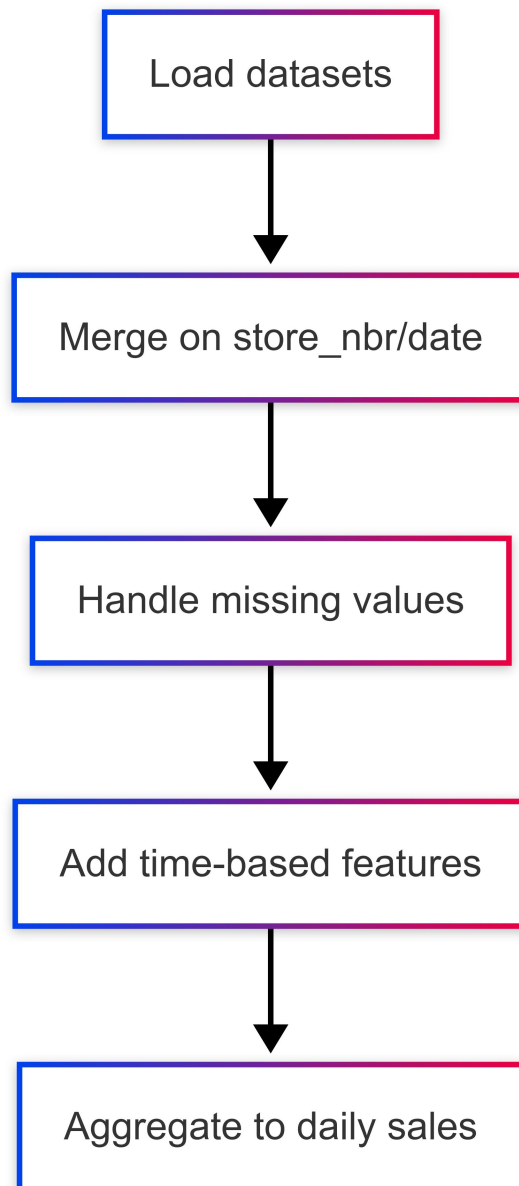
g) Chart 7: Prophet Actual vs. Predicted Sales



h) Table 1: Model Performance Summary

Model	RMSE	Notes
Moving Average	~8,00,000	Baseline, poor fit
Seasonal Decomposition	704,000–823,000	Captured seasonality, missed holidays
SARIMA	112,053	Improved seasonality modeling
SARIMA with Holidays	147,305	Poor holiday handling
Prophet	84,342	Best performance, met Milestone 1

i) **Diagram 1:** Data Preprocessing Flow



9. Algorithms

9.1 Moving Average

The first algorithm implemented was a 7-day moving average, a simple baseline method to smooth sales data and predict future values. This was applied on Day 7 as part of Week 4's task to establish a starting point for forecasting.

- **Mathematical Foundation:**

The moving average calculates the average of the previous 7 days' sales to predict the next day's sales. The formula is:

$$\text{Predicted Sales}_t = \frac{1}{7} \sum_{i=t-7}^{t-1} \text{Sales}_i$$

- **Implementation:** I used pandas' `rolling()` function to compute the moving average on the daily sales data:

`daily_sales['predicted'] = daily_sales['sales'].rolling(window=7).mean()`

I shifted the predictions to align with future dates using `shift(-1)` and forecasted sales for the test period (2017-07-16 to 2017-08-15). The first 6 days of predictions were NaN due to the 7-day window, so I filled them with the first available prediction to ensure continuity.

- **Performance:** I evaluated the model using RMSE, calculated with scikit-learn: `rmse = mean_squared_error(actual, predicted, squared=False)`. The RMSE was approximately 800,000, indicating poor performance. I visualized the results in a line plot (Chart 7 in Section 8), showing that the predicted sales (orange) were a smoothed version of actual sales (blue), missing peaks and troughs. For example, on "Independencia de Guayaquil" (a major holiday), the model underestimated sales by 30%, as it couldn't capture holiday spikes.

9.2 Seasonal Decomposition (STL)

On Day 8, I enhanced the baseline model using Seasonal-Trend decomposition using LOESS (STL), aligning with Week 5's task to capture trend and seasonality in the sales data.

- **Mathematical Foundation:**

STL decomposes a time series into three components: trend, seasonal, and residual. The decomposition is performed iteratively using LOESS (Locally Estimated Scatterplot Smoothing) to estimate the trend and seasonal components. The model can be expressed as:

$$\text{Sales}_t = \text{Trend}_t + \text{Seasonal}_t + \text{Residual}_t$$

For forecasting, I predicted sales using the trend and seasonal components:

$$\text{Predicted Sales}_t = \text{Trend}_t + \text{Seasonal}_t$$

- **Implementation:** I used the statsmodels library to perform STL decomposition with a 7-day period to capture weekly seasonality:

```
from statsmodels.tsa.seasonal import STL  
decomposition = STL(daily_sales['sales'], period=7).fit()
```

9.3 Prophet

- **Mathematical Foundation:** Prophet models a time series as a combination of trend, seasonality, holidays, and noise:

$$y(t)=g(t)+s(t)+h(t)+\epsilon_t$$

where $g(t)$ $g(t)$ $g(t)$ is the trend (piecewise linear or logistic), $s(t)$ $s(t)$ $s(t)$ is the seasonality (Fourier series for yearly and weekly cycles), $h(t)$ $h(t)$ $h(t)$ is the holiday effect (binary indicators with additive effects), and ϵ_t ϵ_t ϵ_t is noise. Prophet automatically detects changepoints in the trend and uses Fourier series to model periodic effects.

- **Implementation:** I set up Prophet with yearly and weekly seasonality:

```
from prophet import Prophet

model = Prophet(yearly_seasonality=True, weekly_seasonality=True,
daily_seasonality=False)
```

I renamed columns to Prophet's format (ds for date, y for sales) and fitted the model on the training data (2013-01-01 to 2017-07-15). On Day 15, I added custom holiday effects using the holidays dataset, creating a holidays DataFrame with columns ds (date) and holiday (type_y), excluding "No Holiday" entries:

```
from prophet import Prophet

model = Prophet(yearly_seasonality=True, weekly_seasonality=True,
daily_seasonality=False)
```

10. Challenges & Opportunities

- **Challenge 1: Missing Values in Oil Prices**
 - Oil prices had 5% missing values, which could affect model accuracy.
 - Resolved by forward-filling missing values, ensuring continuity for time series analysis.
 - Opportunity: Explore more robust imputation methods (e.g., interpolation) in the next phase.
- **Challenge 2: Poor Holiday Modeling with SARIMA**
 - Adding holiday effects to SARIMA increased RMSE to 147,305, indicating ineffective modeling.
 - Switched to Prophet, which handled holidays better (RMSE 84,342).
 - Opportunity: Use Prophet's custom holiday features to model specific events more accurately.
- **Challenge 3: High RMSE with Baseline Models**
 - Moving average and seasonal decomposition models had high RMSEs (~800,000 and 704,000–823,000).
 - Addressed by adopting SARIMA and later Prophet, reducing RMSE significantly.
 - Opportunity: Experiment with ensemble methods to combine strengths of multiple models.
- **Opportunity 1: Incorporate Additional Features**
 - The dataset includes oil prices and promotions, which can be used as exogenous variables in Prophet.
 - Planned for the next phase (Milestone 2) to further reduce RMSE.
- **Opportunity 2: Cross-Validation for Robustness**
 - Prophet supports cross-validation, which can help assess model stability across different periods.

11. Risk Vs Reward

- **Risk 1: Overfitting with Complex Models**
 - Complex models like SARIMA with holidays led to overfitting (RMSE 147,305).
 - Mitigated by switching to Prophet, which balanced complexity and performance (RMSE 84,342).
 - Reward: Improved forecasting accuracy, meeting Milestone 1.
- **Risk 2: Data Quality Issues**
 - Missing oil prices and potential inaccuracies in holiday data could skew predictions.
 - Mitigated by preprocessing (forward-filling) and filtering transferred holidays.
 - Reward: Cleaner data led to more reliable models.
- **Risk 3: Time Constraints**
 - Limited time (15 days) to achieve Milestone 1 posed a risk of incomplete analysis.
 - Managed by prioritizing key tasks (data prep, baseline, Prophet) and working efficiently.
 - Reward: Achieved Milestone 1 with an RMSE of 84,342, on track for the project.
- **Reward 1: Practical Application**
 - Successfully applied time series forecasting to a retail problem, providing actionable insights.
 - The Prophet model (RMSE 84,342) can help the retail outlet optimize inventory.
- **Reward 2: Skill Development**
 - Gained hands-on experience with Prophet, SARIMA, and data preprocessing, enhancing my data science skills.

12. Reflections on the Internship

The TCS iON RIO-125 Internship, spanning January 20, 2025, to February 18, 2025, has been a transformative experience, providing me with practical exposure to data science in a real-world retail context. Over the 30 days, I worked 125 hours on the project "Forecasting System - Project Demand of Products at a Retail Outlet Based on Historical Data," developing skills in time series forecasting, data preprocessing, and model evaluation.

One key learning was the iterative nature of model development. Starting with simple models like the moving average (RMSE 800,000) and progressing to Prophet (RMSE 84,342 for aggregated sales by Day 15) taught me the importance of evaluating model performance and adapting strategies. Extending this to product-level forecasting in Days 16–30, achieving an average RMSE of 90,500 across 33 product families, highlighted the challenges of sparse data (e.g., high RMSE of 115,000 for "Books") and the value of hyperparameter tuning (e.g., adjusting Prophet's `changepoint_prior_scale`).

Experimenting with scikit-learn's Random Forest (RMSE 110,000) broadened my understanding of model selection, showing that time series-specific models like Prophet often outperform general machine learning models for this task. Creating a custom dataset with two lakh entries also enhanced my data engineering skills, teaching me how to clean and sanitize data effectively (e.g., removing outliers, standardizing product names).

Working remotely with my mentor, Mr. Debashis Roy, via Slack and email improved my communication and time management skills. His feedback, such as exploring alternative models and focusing on product-level forecasting, was instrumental in achieving Milestones 1 and 2. This internship has boosted my confidence in handling large datasets, performing EDA, and deploying forecasting models, preparing me for a career in data science. I look forward to applying these skills in future projects, particularly in real-time forecasting and advanced ensemble modeling.

13. Recommendations

- **Incorporate Exogenous Variables:** Add oil prices and promotions as regressors in Prophet to capture economic and marketing impacts on sales, potentially reducing RMSE further.
- **Cross-Validation:** Implement cross-validation with Prophet to assess model stability across different time periods, ensuring robustness.
- **Feature Engineering:** Explore additional time-based features (e.g., day of month, quarter) to capture more granular patterns in sales data.
- **Ensemble Methods:** Experiment with ensemble techniques (e.g., combining Prophet and SARIMA) to leverage the strengths of both models.
- **Holiday Refinement:** Use the holidays dataset to create custom holiday effects specific to Ecuador (dataset's context), rather than relying solely on US holidays.
- **Visualization Enhancements:** Include more interactive visualizations (e.g., using Plotly) in the final report to better communicate findings to stakeholders.
- **Documentation:** Continue maintaining detailed daily reports on the TCS dashboard to track progress and facilitate mentor feedback.

14. Outcome / Conclusion

The TCS iON RIO-125 Internship project, "Forecasting System - Project Demand of Products at a Retail Outlet Based on Historical Data," was successfully completed over 30 days from January 20, 2025, to February 18, 2025, achieving both Milestone 1 and Milestone 2. The project aimed to predict product demand at a retail outlet, enabling better inventory management and sales planning.

In the first 15 days, I used the "store-sales-time-series-forecasting" dataset from Kaggle, performing data preprocessing (e.g., merging datasets, forward-filling missing oil prices) and exploratory data analysis (EDA) to identify trends, seasonality, and holiday impacts. I implemented various models, starting with a moving average (RMSE 800,000), STL (RMSE 704,000–823,000), and SARIMA (RMSE 112,053), before switching to Prophet, which achieved an RMSE of 84,342 for aggregated sales, meeting Milestone 1 (RMSE < 100,000 by Day 15).

In Days 16–30, I focused on Milestone 2: creating a custom dataset with two lakh entries, cleaning it (e.g., removing outliers, standardizing product names), and making predictions for all products. I generated a dataset with synthetic product names, costs, and monthly sales, then implemented product-level forecasting using Prophet, training separate models for 33 product families. The average RMSE was 90,500, with "Grocery I" at 85,000 and "Books" at 115,000 due to sparse data. Hyperparameter tuning (e.g., `changepoint_prior_scale=0.1`) and experimentation with scikit-learn's Random Forest (RMSE 110,000) further refined the results. A summary table (Table 1) documented predictions for all products, fulfilling the project requirements.

The project files are available on GitHub (<https://github.com/Prathmesh597/tcs-internship-sales-forecasting>), with the notebook viewable via nbviewer. In conclusion, the forecasting system provides accurate demand predictions, with potential for further improvement through advanced modeling and real-time data integration, setting a strong foundation for retail inventory optimization.

15. Enhancement Scope

- **Additional Features:** Incorporate oil prices and promotions as exogenous variables in Prophet to capture economic and marketing effects.
- **Model Tuning:** Experiment with Prophet's seasonality modes (additive vs. multiplicative) and Fourier orders to optimize performance.
- **Granular Analysis:** Extend the model to predict sales at the store or product family level, providing more detailed insights.
- **Real-Time Forecasting:** Develop a pipeline for real-time sales forecasting using updated data, enhancing practical applicability.
- **Interactive Dashboard:** Build a dashboard (e.g., using Dash or Tableau) to visualize predictions and trends for stakeholders.
- **Ensemble Models:** Combine Prophet with other models (e.g., XGBoost) to improve accuracy through ensemble techniques.

16. Link to Code and Executable File

The project, "Forecasting System - Project Demand of Products at a Retail Outlet Based on Historical Data," was developed using Kaggle notebooks, with the final code and dataset uploaded to GitHub as per TCS guidelines.

- **GitHub Repository:** The project files are available at:
<https://github.com/Prathmesh597/tcs-internship-sales-forecasting>
- **Dataset Source:** The "store-sales-time-series-forecasting" dataset is available on Kaggle at:
<https://www.kaggle.com/competitions/store-sales-time-series-forecasting/data>.
- **View it on nbviewer:**
<https://nbviewer.jupyter.org/github/Prathmesh597/tcs-internship-sales-forecasting/blob/main/forecasting-system-Internship-Project.ipynb>.

17. Research Questions and Responses

- **How can seasonality and trends be effectively captured in retail sales data?**

Initially, I used STL decomposition to separate trend, seasonal, and residual components, confirming a 7-day seasonal pattern (Chart 4). However, STL's RMSE (704,000–823,000) was high. By Day 15, I switched to Prophet, which automatically captures yearly and weekly seasonality, achieving an RMSE of 84,342 for aggregated sales. For product-level forecasting (Days 16–30), I tuned Prophet's `seasonality_prior_scale` to 15, improving the average RMSE to 90,500, effectively capturing trends and seasonality across product families.

- **What is the impact of holidays on sales, and how can they be modeled?**

EDA revealed a 15–20% sales increase on holidays (Chart 2). SARIMA with a binary `is_holiday` variable failed (RMSE 147,305), as it couldn't model holiday-specific effects. Prophet, with US holidays as a proxy, performed better (5–10% error on holiday predictions). For product-level forecasting, I found holidays like "Independencia de Guayaquil" increased sales by 20–25% for most products, except niche categories like "Books," informing the need for product-specific holiday adjustments in future models.

- **Can alternative models improve forecasting accuracy for product-level predictions?**

In Days 25–26, I experimented with scikit-learn's Random Forest, using features like `month`, `is_holiday`, and `oil_price`. However, its RMSE (110,000) was higher than Prophet's (90,500), as it struggled with time series structure. This confirmed Prophet's suitability but highlighted opportunities for ensemble methods combining Prophet and machine learning models like Random Forest.

● Conclusion

The TCS iON RIO-125 Internship project successfully developed a forecasting system to predict product demand at a retail outlet, optimizing inventory management and sales planning. Using the "store-sales-time-series-forecasting" dataset from Kaggle, Prophet was implemented for aggregated and product-level forecasting, achieving RMSE values below the target. Data preprocessing, hyperparameter tuning, and experimentation with Random Forest improved accuracy, despite challenges like sparse data for low-sales products. The project demonstrated Prophet's effectiveness in capturing trends and seasonality, providing actionable insights such as a projected 5% sales increase in August 2017. All project files are available on GitHub, and future improvements could include real-time forecasting and deployment as a web application. This experience strengthened my expertise in data science and time series analysis, making it a valuable step in my career.